

Perspective

Lessons from complex systems science for AI governance

Noam Kolt,^{1,2,*} Michal Shur-Ofry,¹ and Reuven Cohen³

¹Faculty of Law, Hebrew University, Jerusalem, Israel

²School of Computer Science and Engineering, Hebrew University, Jerusalem, Israel

³Department of Mathematics, Bar-Ilan University, Ramat Gan, Israel

*Correspondence: noam.kolt@mail.huji.ac.il

<https://doi.org/10.1016/j.patter.2025.101341>

THE BIGGER PICTURE Complex adaptive systems are systems composed of numerous interacting agents that self-organize and adapt to changing environments. Artificial intelligence (AI) systems increasingly exhibit features that are characteristic of complex adaptive systems. These features make it difficult to reliably predict and steer the impact of AI systems on society. Traditional governance approaches, which often assume linear cause-and-effect relationships, are not up to the task. By leveraging insights from complex systems science, we propose governance principles tailored to the reality of modern AI. These principles include early and scalable intervention, adaptive institutional design, and risk thresholds that reflect the nonlinearity of AI systems and their interactions with other sociotechnical structures. Reframing AI governance through the lens of complexity can help law and policy keep pace with the rapid changes arising from this consequential technology.

SUMMARY

The study of complex adaptive systems, pioneered in physics, biology, and the social sciences, offers important lessons for artificial intelligence (AI) governance. Contemporary AI systems and the environments in which they operate exhibit many of the properties characteristic of complex systems, including nonlinear growth patterns, emergent phenomena, and cascading effects that can lead to catastrophic failures. Complex systems science can help illuminate the features of AI that pose central challenges for policymakers, such as feedback loops induced by training AI models on synthetic data and the interconnectedness between AI systems and critical infrastructure. Drawing on insights from other domains shaped by complex systems, including public health and climate change, we examine how efforts to govern AI are marked by deep uncertainty. To contend with this challenge, we propose three desiderata for designing a set of complexity-compatible AI governance principles comprised of early and scalable intervention, adaptive institutional design, and risk thresholds calibrated to trigger timely and effective regulatory responses.

INTRODUCTION

Discussion of the impact of artificial intelligence (AI) and approaches to governing the technology have become increasingly polarized. Scholars and practitioners concerned about the risks of AI systems fiercely debate the appropriate goals, scope, and timing of regulatory policy and intervention,^{1,2} often divided along disciplinary lines or research communities.³ The discourse is, to a large extent, influenced by conceptual framing. Some characterize AI as a highly consequential software product or service,^{4–6} while others characterize AI as a societal-scale transformation that presents unprecedented risks.^{7,8}

We propose a different lens, grounded in decades of interdisciplinary research in physics, biology, and the social sciences: analyzing AI systems, their development process, and the envi-

ronments in which they operate as complex systems. Complex systems are systems comprising multiple interacting components. Examples of such systems, in the natural and human world, include insect colonies, urban environments, social networks, and financial markets.^{9–12}

Complex systems science demonstrates that complex systems of different kinds share common traits, including the emergence of systems-level properties and patterns despite the absence of central control or design and nonlinear dynamics that defy simple cause-effect relations. Small changes in a complex system's topology and interactions among its components may result in very different overall effects. Complex systems thus entail inherent unpredictability and are susceptible to rare but substantial cascades and the materialization of catastrophic failures with far-reaching consequences (Box 1).^{9,13–17}



Box 1. Defining complex systems

“[A] complex system ... [is] ... one made up of a large number of parts that interact in a non-simple way. In such systems, the whole is more than the sum of its parts... in the important pragmatic sense that, given the properties of the parts and the laws of their interaction, it is not a trivial matter to infer the properties of the whole.” — Herbert A. Simon, *The Architecture of Complexity*¹⁹⁰

Methodologies developed in complex systems science enable researchers to better understand complex systems and, where appropriate, design policies to address the associated societal challenges. Drawing on studies that suggest that central aspects of contemporary AI systems bear the hallmarks of complexity,^{18–30} we make three primary contributions. First, we unpack the characterization of AI systems as complex systems. Second, we explore the implications of this characterization for the challenges involved in governing AI. Third, drawing on insights from complexity and other domains shaped by complex systems, we propose a series of complexity-compatible principles to assist policymakers in developing more effective mechanisms for governing AI.

AI AND COMPLEXITY

Our characterization of AI systems as complex systems focuses on several properties increasingly identified in AI systems, the processes through which they are trained, and the environments in which they are deployed. The properties we focus on are nonlinear growth, unpredictable scaling and emergence, feedback loops, cascading effects, and susceptibility to catastrophic failures. While the list is not exhaustive and varies substantially across different forms of AI and application domains, these properties illuminate some of the distinctive governance challenges posed by AI.

A note on terminology: while there is no accepted consensus as to the definition of “AI,” we use the terms “AI” and “AI systems” expansively³¹ so as to include a wide range of computational systems with varying levels of autonomy (as defined in the EU AI Act³²), recognizing their embeddedness in sociotechnical environments.^{3,18,19,25–27,29}

Nonlinear growth

In recent years, there has been an exponential increase in many of the key inputs into AI development. The computational resources for training AI models have, on average, grown by a factor of four to five each year between 2010 and 2024.³³ The size of datasets used for training has also increased significantly. For example, language training dataset size has increased by a factor of three each year in recent years.³⁴ Meanwhile, efficiency of computation has continued to improve exponentially,^{35,36} alongside corporate investment that has increased by more than an order of magnitude.³⁷

During this period, the capabilities of AI systems have improved dramatically. AI systems can now outperform humans on some tasks,³⁸ including certain tasks relating to visual question answering, natural language understanding, and text annotation.^{37,39} Several benchmarks previously used to evaluate AI systems have, due to improvements in the capabilities of AI, been rendered obsolete.^{40,41} AI systems have also achieved superhuman feats in various scientific fields, including biology,^{42,43} mathematics,⁴⁴ weather forecasting,⁴⁵ and materials science.⁴⁶

Importantly, as we illustrate, these improvements in performance may themselves exhibit properties of complexity.

Scaling, emergence, and unpredictability

The effect of increases in the inputs into widely used AI systems, especially foundation models, on the performance of those systems resembles patterns characteristic of other complex systems. This phenomenon has been observed both in model training since the advent of foundation models⁴⁷ and in model inference, following the development of reasoning models (i.e., models that “think” using chain of thought at run time).⁴⁸

In model training, cross-entropy loss—the main metric used to measure training performance—has been shown to scale (decrease) in a power-law relationship with model size, dataset size, and the amount of compute used in training.^{49–52} These “scaling laws” suggest that the performance of AI models (measured by cross-entropy loss) may continue to improve with increases in the inputs used in training. However, the specific capabilities acquired by these models in practice, that is, their ability to perform particular real-world tasks, remain highly unpredictable and can appear to emerge suddenly.^{24,53–55} An early illustration of this phenomenon was observed in 2020 with GPT-3, which, although structurally similar to prior models, due to its larger size, gained the qualitatively new ability to learn to perform new tasks after being provided a few demonstrations of those tasks (known as “few-shot learning”).⁵⁶

More recently, progress in the development of reasoning models—such as OpenAI’s o series of models⁵⁷ and DeepSeek’s R series of models⁵⁸—demonstrates an equivalent phenomenon. Increasing the scale of inference (test-time) compute has resulted in systems gaining qualitatively new abilities, including exceeding human PhD-level accuracy on certain STEM-related benchmarks.^{59–63}

Seen through the lens of complex systems science, the sudden emergence of new capabilities in AI models can be analogized to phase transitions in physical and biological systems,^{64,65} such as water freezing or boiling when it reaches a certain temperature or the emergence of cognition from multiple neural interactions.^{66,67} Similarly, AI systems appear to acquire new, qualitatively different abilities when a certain threshold is reached, either in training or at inference. The exact scope and nature of new AI abilities, however, are unpredictable. For instance, in late 2024, OpenAI’s o3 model was able to solve over 25% of the problems in the FrontierMath benchmark, surpassing the 2% achieved by previous models and defying Terence Tao’s prediction that the problems would “resist AIs for several years at least.”⁶⁸

Feedback loops

Like other complex systems, AI systems interact with their environments and are prone to feedback loops that can generate self-reinforcing processes. These can occur, for instance,

when the output of an AI model influences human behavior that is then incorporated into the data used to refine the model or train future models. For example, various algorithms used to predict housing prices can influence real-world housing prices, which then influence future price predictions, and so on.⁶⁹ This kind of feedback loop, in which predictions that support decisions influence the very outcomes they aim to predict, is known as “performative prediction.”^{70–72} Feedback loops also commonly arise in the context of content recommendation. Recommender systems respond to users’ selection of content by recommending similar content, which, in turn, reinforces users’ existing content preferences.^{73–76}

The widespread use of foundation models exacerbates such feedback loops. Because the outputs of foundation models are increasingly incorporated into publicly available data repositories, which are then included in the training data of future models, errors and biases in earlier models could compound with each successive generation of models.^{77–79} For example, anti-consumer biases in language models used to perform legal tasks could intensify if the biased outputs of those models are used to train future models.⁸⁰

Feedback loops might also ensue as humans tasked with annotating data for training AI models outsource their work to other AI models^{81–83} or are influenced by their use of AI models.⁸⁴ In addition, training models on large quantities of synthetic data (i.e., data generated by other AI models)^{85–88} can, in some circumstances, degrade the quality of the resulting models.^{89–94} This may be exacerbated by the fact that detecting AI-generated content (e.g., by using watermarks)⁹⁵ and excluding it from training datasets remain difficult.^{96,97} Studies suggest that a growing fraction of content on the internet is already dominated by synthetic content, including synthetic content produced by models that are themselves trained on synthetic content.^{98–100}

Novel feedback loops could also arise as AI models are increasingly used to evaluate the safety of other AI models^{101–104} or assist in conducting AI safety research.^{105,106} While forecasting the precise contours of these interactions will likely be impossible, suffice it to say that research in complex systems science suggests that these phenomena could lead to rapid and potentially dangerous self-reinforcing processes, especially in the case of interconnected systems and networks.^{107,108}

Interconnectedness, cascading effects, and catastrophic failures

Complex systems science sheds light on the vulnerability of interdependent networks to cascading effects, whereby damage to a small number of nodes (i.e., components comprising the system) in one network can have an outsized impact on other interconnected systems, potentially causing large-scale damage.^{9,13–17} For example, power outages can cause internet outages that then cause further power outages, which can affect additional interconnected networks, such as telecommunications networks.¹⁴

Similar dynamics could—and perhaps already do—arise in AI, especially when AI systems are integrated into other systems. The prevailing AI paradigm in which foundation models perform downstream applications across multiple domains is particularly

vulnerable to cascading effects. Minor defects in foundation models can propagate across the myriad settings in which they are deployed.^{47,109} While the homogenization introduced by foundation models promotes efficiency (expensive-to-train models can be cheaply reused and adapted to many applications), it also gives rise to new risks familiar to complexity researchers. Safety failures resulting from foundation models might not be independent or isolated from one another but rather correlated and connected.¹¹⁰ For example, systems that exhibit misalignment in one context (e.g., they produce insecure code) tend to exhibit misalignment in other, ostensibly unrelated contexts (e.g., they provide malicious advice and act deceptively).¹¹¹ Meanwhile, vulnerability to a particular type of adversarial attack can diffuse across multiple domains in which a model or agent is deployed or across different models or agents built using similar architecture.^{112,113}

Cascading effects could compound as AI systems are integrated into external networks and infrastructure. For example, autonomous agents tasked with pursuing complex goals in safety-critical domains, such as financial markets and essential services, could have highly unpredictable and adverse consequences.^{114–116} As autonomous agents are increasingly integrated into other systems, potential cascading effects could become even broader and harder to predict.¹¹⁷ A central factor in assessing these effects and associated risks is the level of interconnectedness between AI systems and the other systems with which they interact.¹¹⁸ Higher levels of interconnectedness imply more vulnerability of risks percolating from an AI system to other systems.¹¹⁹ Meanwhile, AI systems that operate in closed environments or in settings with only limited interconnectedness likely pose less severe risks.

As a result of interconnectedness, feedback loops, and cascading effects, complex systems are particularly susceptible to catastrophic failures. For instance, interconnected feedback loops can upend financial markets in unpredictable high-impact events sometimes described as black swans.¹²⁰ AI systems could give rise to similar risks.^{7,8} To illustrate, a malfunctioning AI system used to control wastewater treatment facilities might not only cause direct harm by discharging untreated effluent but could also have wider adverse effects on human health and marine life.¹²¹ These catastrophic failures could become more acute if AI systems are integrated into critical infrastructure. For instance, if AI systems are used to control water infrastructure that cools data centers used to train or operate AI systems, then single-system failures—whether resulting from accidental malfunction or malicious adversarial attack—could have far wider consequences.^{121–123}

Importantly, catastrophic failures from AI might not necessarily materialize solely due to the defects in a particular AI system percolating to other systems. Instead, catastrophic failures may arise through the interaction of AI systems with broader sociotechnical structures.^{3,18,19,21,22,25–27,29,30} For example, economic incentives and corporate governance structures may prompt companies to deploy AI systems and enable their use in high-stakes domains without sufficient safeguards.^{53,124–126} The rapid diffusion and adoption of these systems dramatically increase the surface area of potential catastrophic failures and present difficult governance challenges.

LESSONS FOR AI GOVERNANCE

Understanding AI systems as complex systems illuminates important governance challenges, many of which are overlooked by current regulatory frameworks.^{118,125,127} To illustrate, the European Union's AI Act defines "systemic risk" from AI as "actual or reasonably foreseeable negative effects on public health, safety, public security, fundamental rights, or the society as a whole"³² (Art. 3(65)). This definition arguably equates "systemic risk" with "large-scale harm." By drawing this parallel and grounding the definition of "systemic risk" in "actual or reasonably foreseeable negative effects," the EU AI Act overlooks the unpredictable and cascading nature of risks in interconnected systems,¹²⁸ the interactions between those systems and the environments in which they operate,¹²⁹ and the relationship between systemic risk and conventional risks.¹³⁰ Moreover, despite referring to "systemic risk" and purporting to address this class of risk in its operative provisions³² (Art. 55), the Act does not, in fact, adopt a "systems thinking" approach to AI governance, which would entail, among other things, more robust engagement with complex systems science.¹³¹ Employing perspectives from complex systems science would allow policymakers to draw on regulatory insights regarding complex systems in other domains, including climate policy,^{132–134} financial regulation,^{135–137} and public health,¹³⁸ and thereby develop better calibrated guiding principles for tackling the governance challenges posed by AI.

Regulating under deep uncertainty

While the regulation of any moving target is difficult,^{139–142} the regulation of AI systems characterized by rapid development, emergent properties, feedback loops, and unpredictable cascading effects is a particularly thorny problem.^{125,127,143} Neither technologists nor policymakers can reliably predict the capabilities of AI systems or accurately forecast their negative externalities.^{53,125} Regulatory efforts, whether targeted at model development or deployed systems and applications, must contend with an ongoing information deficit^{144–146} and deep uncertainty.^{130,147,148} The problem, at its core, is that by the time the capabilities and real-world ramifications of AI systems are properly understood, it may be too late to intervene effectively—a challenge familiar to policymakers in other domains.¹³⁹

In light of these challenges, we propose three desiderata for designing AI governance mechanisms: (1) policymakers should have the capacity and resources to take early and scalable regulatory action, (2) regulatory action should be dynamic and highly responsive to changing conditions, and (3) policymakers should adopt complexity-compatible risk thresholds with respect to AI systems that exhibit properties characteristic of complex systems.

Early and scalable intervention

When risks cascade in complex systems, policymakers must be able to respond early, rapidly, and at scale.^{115,125} For example, to prevent large-scale economic harm from vulnerabilities in a widely used automated stock-trading tool, regulators may need to intervene before the harm has (fully) materialized. Counterintuitively, the case for robust intervention to govern complex systems, including AI technologies, may decline over time.¹³⁸ While early intervention (made on the basis of only limited infor-

mation) could prevent the relevant harm, intervention at a later point (made on the basis of more complete information) may in fact no longer be effective.¹³⁸ By analogy, lockdowns and border closures designed to prevent the spread of a pandemic are far more effective and hence more justifiable earlier in time (despite the absence of complete information), before the pandemic has spread beyond the ability to contain it, after which such mandates may no longer be as effective.¹³⁸ A similar dynamic could apply to AI technologies that exhibit emergent properties, diffuse rapidly and nonlinearly, and lead to cascading effects that could cause large-scale harm.

Apart from the timing of intervention, mechanisms for governing AI systems must also operate at sufficient scale.^{125,134,149} Continuing with the example of vulnerabilities in a widely used automated stock-trading tool, for governance mechanisms to be effective, they will need to operate successfully across a very large number of actors, institutions, and environments that interact with the tool in question. Consequently, certain conventional governance mechanisms, such as manual human oversight and evaluation, may be ineffective,¹⁵⁰ while more scalable mechanisms, such as automated oversight and evaluation, may, despite their shortcomings and potential risks, become necessary.^{103–105}

Adaptive governance

Even if the above desideratum is met, governance institutions will nonetheless need to adapt to new conditions arising due to hard-to-predict changes in AI systems, their usage, and the broader sociotechnical context in which they operate.^{151,152} To this end, policymakers should draw on the principles of adaptive management and resilience proposed in the field of climate policy and environmental governance.^{153–155} According to these principles, governance mechanisms that aim to regulate complex systems should not be static institutions but rather feedback-driven processes that iteratively respond and adapt to new information while preserving overarching societal goals and values.^{130,134,144,156} This dynamic approach to governance is especially crucial for mitigating cascading failures of complex systems.^{157–160}

As illustrated in Table 1, prominent AI governance frameworks include mechanisms for adaptation and change. Notably, these mechanisms for adaptation do not specify the type of information required to bring about changes in the corresponding governance framework. For example, it is unclear what information the EU would need to receive in order to add or remove AI systems or applications from the list of high-risk systems in the EU AI Act.³² That being said, this feature of regulatory frameworks is not necessarily a defect. Perspectives from complex systems suggest that overly fine-grained rules that attempt to anticipate every possible contingency are inherently limited.¹⁶¹ Accordingly, the use of open-ended standards in AI regulation that accommodate regulatory discretion and responsiveness—without stipulating the precise type of information required to trigger regulatory action—has notable advantages and may, on balance, be preferable to more prescriptive approaches.

Nevertheless, to be effective, adaptive governance must be guided by up-to-date information concerning the systems being governed.^{144–146,168–170} In the case of AI, policymakers must continually acquire information about the current and anticipated capabilities, trends, and impacts of AI systems.^{171–174} In other

Table 1. Adaptation mechanisms in prominent governance frameworks

	Type of framework	Adaptation mechanisms
US National Institute of Standards and Technology (NIST) AI Risk Management Framework (January 2023) ¹⁶²	non-binding practice and policy framework	describes the framework as a “living document” and refers to current document as v.1.0; stipulates that NIST will regularly review the document, including with formal public input
China Interim Measures for the Management of Generative Artificial Intelligence Services (August 2023) ¹⁶³	binding obligations on generative AI service providers	these interim measures are likely to be superseded by the draft Artificial Intelligence Law of the People’s Republic of China first circulated in March 2024 ¹⁶⁴
US Executive Order on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence (October 2023) ¹⁶⁵	binding reporting requirements and mandatory government actions	requirements carried out by various government agencies that exercise significant discretion; like other US executive orders, the order can be modified or revoked by the President—and was revoked by President Trump in January 2025 ¹⁶⁶
European Union Artificial Intelligence Act (August 2024) ³²	binding cross-sector regulation and establishment of new regulatory institutions	while the Act itself will be difficult to amend, it includes mechanisms for amending certain key provisions and further changes through implementing acts, delegated acts, externally determined standards, and a code of practice for general-purpose AI ¹⁶⁷

words, they must engage in “evidence-seeking” policy.^{175,176} Adaptive governance also implies that policymakers should be cognizant of potential abrupt changes in a system’s performance or behavior and of the possibility that such changes will quickly percolate and affect interconnected systems.

Current regulatory frameworks have made significant progress in tackling this information problem, establishing multiple mechanisms for furnishing policymakers with decision-relevant information. For instance, the EU AI Act requires companies to keep detailed records of certain “high-risk” AI systems and proposes mechanisms for monitoring these systems and reporting safety incidents³² (Arts. 11–12, 72–73). Notwithstanding these mechanisms, deciding how to interpret the information gathered and whether (or how) to act upon it still presents a significant challenge for policymakers.

Complexity-compatible risk thresholds

What threshold of risk from AI should trigger regulatory intervention?¹⁷⁷ What information would constitute sufficient evidence that such a threshold has been reached? Regulators often dodge these questions by postponing governance decisions until the relevant “evidentiary burden” is satisfied.¹⁷⁸ The problem with this approach is that, because many AI systems exhibit properties characteristic of complex systems, such information may only become available at a time after which intervention has become more costly or less effective.^{138,139,141,179,180} (This approach can also be exploited by interested parties seeking to delay or obstruct effective regulation.^{175,176})

Consequently, to intervene effectively, policymakers may need to relax the policy-relevant informational threshold and resort to “satisficing”¹⁸¹—i.e., making governance decisions on the basis of incomplete information collected at an earlier stage in the technology’s development and use.¹³⁸ For example, policymakers may need to amend AI safety standards upon receiving interim red-teaming results that indicate certain dangerous capabilities prior to receiving the final results, let alone comprehensive studies establishing the precise probability or magnitude of the relevant risks. In such cases, rather than wait until more detailed or complete information is avail-

able, regulators should familiarize themselves with the patterns characteristic of complex systems in order to evaluate the potential risks from AI systems and design appropriate interventions.

One potential route is to employ a legal doctrine known as the “precautionary principle.” The principle, which is used in environmental governance and public health policy, supports pre-emptive regulatory intervention before harms have (fully) materialized or risks are established conclusively, often requiring actors interested in pursuing a potentially risky activity to first prove its safety.^{182,183} A prominent criticism of the precautionary principle—which, in the case of AI, may require robust technical safety guarantees^{184,185} or other forms of assurance^{186,187}—is that it does not withstand cost-benefit analysis, i.e., it unduly limits or forgoes the gains from new technology.^{179,188,189}

However, as explored in the context of pandemic responses, insights from complexity can help refine and calibrate the precautionary principle. In particular, regulators should consider whether the relevant risks will likely spread swiftly and exponentially and thereby pose a grave systemic risk.¹³⁸ When this is the case, the costs of postponing regulatory intervention until more complete information is obtained are often multiplicative, such that delay can be orders of magnitude costlier than early intervention. For example, refraining from intervening to prevent failures in an AI system connected to critical infrastructure could result in costly damage that rapidly percolates into other safety-critical systems.^{14–17,158,159} Conversely, the costs of early intervention (e.g., requiring additional guardrails in response to interim red-teaming results) are often linear and additive. Seen through the lens of complexity, cost-benefit analysis can, in certain cases, support a precautionary approach to governing AI. A central consideration in this analysis is the level of interconnectedness between AI systems and other sociotechnical systems,¹¹⁸ as well as the potential for feedback loops and cascading effects. Future work will need to examine these considerations in specific settings in order to implement complexity-compatible risk thresholds in practice.

OUTLOOK

The hallmarks of complex adaptive systems increasingly exhibited by AI—nonlinearity, emergence, feedback loops, cascading effects, and the potential for catastrophic failures—underscore the difficult governance challenges facing policymakers. Studying AI through the lens of complexity can guide policymakers to focus on the well-studied patterns of complex systems that are likely to arise in the context of AI systems. Complex systems science helps identify and characterize new risks from AI technologies and points toward more appropriate governance mechanisms. Policymakers addressing the challenges from AI should draw on approaches developed in other domains that confront complexity-related challenges, including climate policy and public health. As AI technologies continue to advance and diffuse, the time is ripe to deepen these interdisciplinary connections.

DECLARATION OF INTERESTS

The authors declare no competing interests.

REFERENCES

- Price, H., and Connelly, M. (2023). AI governance must deal with long-term risks as well. *Nature* 622, 31. <https://doi.org/10.1038/d41586-023-03117-z>.
- Editorial. (2023). Stop talking about tomorrow's AI doomsday when AI poses risks today. *Nature* 618, 885–886. <https://doi.org/10.1038/d41586-023-02094-7>.
- Lazar, S., and Nelson, A. (2023). AI safety on whose terms? *Science* 381, 138. <https://doi.org/10.1126/science.adf8982>.
- Bommasani, R., Kapoor, S., Klyman, K., Longpre, S., Ramaswami, A., Zhang, D., Schaake, M., Ho, D.E., Narayanan, A., and Liang, P. (2024). Considerations for governing open foundation models. *Science* 386, 151–153. <https://doi.org/10.1126/science.adp1848>.
- Kapoor, S., Bommasani, R., Klyman, K., Longpre, S., Ramaswami, A., Cihon, P., Hopkins, A., Bankston, K., Biderman, S., Bogen, M., et al. (2024). Position: on the societal impact of open foundation models. In *Forty-first International Conference on Machine Learning. ICML'24 (JMLR)*.
- Narayanan, A., and Kapoor, S. (2025). AI as normal technology. *Knight First Amend. Inst.*
- Anwar, U., Saparov, A., Rando, J., Paleka, D., Turpin, M., Hase, P., Lubana, E.S., Jenner, E., Casper, S., Sourbut, O., et al. (2024). Foundational challenges in assuring alignment and safety of large language models. *Transactions on Machine Learning Research*. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2404.09932>.
- Bengio, Y., Hinton, G., Yao, A., Song, D., Abbeel, P., Darrell, T., Harari, Y. N., Zhang, Y.Q., Xue, L., Shalev-Shwartz, S., et al. (2024). Managing extreme AI risks amid rapid progress. *Science* 384, 842–845. <https://doi.org/10.1126/science.adn0117>.
- Cohen, R., and Havlin, S. (2010). *Complex Networks: structure, robustness, and function* (Cambridge University Press). <https://doi.org/10.1017/CBO9780511780356>.
- Miller, J.H., and Page, S.E. (2007). *Complex Adaptive Systems: An Introduction to Computational Models of Social Life*. Princeton Studies in Complexity (Princeton University Press).
- Mitchell, M. (2009). *Complexity: A Guided Tour* (Oxford University Press).
- Siegenfeld, A.F., and Bar-Yam, Y. (2020). An introduction to complex systems science and its applications. *Complexity* 2020, 1–16.
- Bashan, A., Berezin, Y., Buldyrev, S.V., and Havlin, S. (2013). The extreme vulnerability of interdependent spatially embedded networks. *Nat. Phys.* 9, 667–672. <https://doi.org/10.1038/nphys2727>.
- Buldyrev, S.V., Parshani, R., Paul, G., Stanley, H.E., and Havlin, S. (2010). Catastrophic cascade of failures in interdependent networks. *Nature* 464, 1025–1028. <https://doi.org/10.1038/nature08932>.
- Gao, J., Bashan, A., Shekhtman, L., and Havlin, S. (2022). *Introduction to Networks of Networks*. In IOP Ebooks Series, 1 (Institute of Physics Publishing).
- Li, W., Bashan, A., Buldyrev, S.V., Stanley, H.E., and Havlin, S. (2012). Cascading failures in interdependent lattice networks: The critical role of the length of dependency links. *Phys. Rev. Lett.* 108, 228702. <https://doi.org/10.1103/PhysRevLett.108.22.8702>.
- Yang, Y., Nishikawa, T., and Motter, A.E. (2017). Small vulnerable sets determine large network cascades in power grids. *Science* 358, eaan3184. <https://doi.org/10.1126/science.aan3184>.
- Dobbe, R. (2022). System safety and artificial intelligence. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency. FAccT '22 (Association for Computing Machinery)*, p. 1584. <https://doi.org/10.1145/3531146.3533215>.
- Hendrycks, D. (2025). *Introduction to AI Safety, Ethics, and Society* (Taylor & Francis).
- Holtzman, A., West, P., and Zettlemoyer, L. (2023). Generative models as a complex systems science: How can we make sense of large language model behavior?. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2308.00189>.
- Leveson, N.G. (2012). *Engineering a Safer World: Systems Thinking Applied to Safety*. Engineering Systems (The MIT Press).
- Macrae, C. (2022). Learning from the failure of autonomous and intelligent systems: Accidents, safety, and sociotechnical sources of risk. *Risk Anal.* 42, 1999–2025. <https://doi.org/10.1111/risa.13850>.
- Nanda, N., Chan, L., Lieberum, T., Smith, J., and Steinhardt, J. (2022). Progress measures for grokking via mechanistic interpretability. In *The Eleventh International Conference on Learning Representations*. <https://openreview.net/forum?id=9XFSbDPmdW>.
- Power, A., Burda, Y., Edwards, H., Babuschkin, I., and Misra, V. (2022). Grokking: Generalization beyond overfitting on small algorithmic datasets. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2201.02177>.
- Rakova, B., and Dobbe, R. (2023). Algorithms as social-ecological-technological systems: an environmental justice lens on algorithmic audits. In *Proceedings of the 2023 ACM Conference on Fairness, Accountability, and Transparency. FAccT '23 (Association for Computing Machinery)*, p. 491. <https://doi.org/10.1145/3593013.3594014>.
- Rismani, S., Shelby, R., Smart, A., Jatho, E., Kroll, J., Moon, A., and Rostamzadeh, N. (2023). From plane crashes to algorithmic harm: Applicability of safety engineering frameworks for responsible ml. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems. CHI '23 (Association for Computing Machinery)*, pp. 1–18. <https://doi.org/10.1145/3544548.3581407>.
- Shelby, R., Rismani, S., Henne, K., Moon, A., Rostamzadeh, N., Nicholas, P., Yilla-Akbari, N., Gallegos, J., Smart, A., Garcia, E., and Virk, G. (2023). Sociotechnical harms of algorithmic systems: Scoping a taxonomy for harm reduction. In *Proceedings of the 2023 AAAI/ACM Conference on AI, Ethics, and Society. AIES '23 (Association for Computing Machinery)*, pp. 723–741. <https://doi.org/10.1145/3600211.3604673>.
- Wei, J., Tay, Y., Bommasani, R., Raffel, C., Zoph, B., Borgeaud, S., Yogatama, D., Bosma, M., Zhou, D., Metzler, D., et al. (2022). Emergent abilities of large language models. In *Transactions on Machine Learning Research*.
- Weidinger, L., Rauh, M., Marchal, N., Manzini, A., Hendricks, L.A., Mateos-Garcia, J., Bergman, S., Kay, J., Griffin, C., Bariach, B., et al. (2023). Sociotechnical safety evaluation of generative AI systems. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2310.11986>.
- Yoo, C. (2024). Beyond algorithmic disclosure for generative AI. *Columbia Sci. Technol. Law Rev.* 25. <https://doi.org/10.52214/str.v25i2.12766>.

31. Schuett, J. (2023). Defining the scope of AI regulations. *Law, Innovation and Technology* 15, 60–82.
32. European Union (2024). Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence and amending Regulations (EC) No 300/2008, (EU) No 167/2013, (EU) No 168/2013, (EU) 2018/858, (EU) 2018/1139 and (EU) 2019/2144 and Directives 2014/90/EU, (EU) 2016/797 and (EU) 2020/1828 (Artificial Intelligence Act) (Text with EEA relevance).
33. Sevilla, J. (2024). Training compute of frontier AI models grows by 4–5x per year. Epoch AI. <https://epoch.ai/blog/training-compute-of-frontier-ai-models-grows-by-4-5x-per-year>.
34. Epoch AI (2025). Machine learning trends. Epoch AI. <https://epoch.ai/trends>.
35. Pilz, K., Heim, L., and Brown, N. (2023). Increased compute efficiency and the diffusion of AI capabilities. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2311.15377>.
36. Ho, A., Besiroglu, T., Erdil, E., Owen, D., Rahman, R., Guo, Z.C., Atkinson, D., Thompson, N., and Sevilla, J. (2024). Algorithmic progress in language models. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2403.05812>.
37. Stanford University (2024). The AI Index 2024 Annual Report. (AI Index Steering Committee, Institute for Human-Centered AI, Stanford University). <https://aiindex.stanford.edu/report/>.
38. Morris, M.R., Sohl-Dickstein, J., Fiedel, N., Warkentin, T., Dafoe, A., Faust, A., Farabet, C., and Legg, S. (2024). Position: Levels of agi for operationalizing progress on the path to agi. In *Forty-first International Conference on Machine Learning*.
39. Gilardi, F., Alizadeh, M., and Kubli, M. (2023). Chatgpt outperforms crowd workers for text-annotation tasks. e2305016120. *Proc. Natl. Acad. Sci. USA* 120. <https://doi.org/10.1073/pnas.2305016120>.
40. Deng, J., Dong, W., Socher, R., Li, L.J., Li, K., and Fei-Fei, L. (2009). Imagenet: A large-scale hierarchical image database. In *2009 IEEE Conference on Computer Vision and Pattern Recognition*, pp. 248–255. <https://doi.org/10.1109/CVPR.2009.5206848>.
41. Wang, A., Pruksachatkun, Y., Nangia, N., Singh, A., Michael, J., Hill, F., Levy, O., and Bowman, S.R. (2019). Superglue: a stickier benchmark for general-purpose language understanding systems. In *Proceedings of the 33rd International Conference on Neural Information Processing Systems (Curran Associates Inc.)*, pp. 3266–3280.
42. Abramson, J., Adler, J., Dunger, J., Evans, R., Green, T., Pritzel, A., Ronneberger, O., Willmore, L., Ballard, A.J., Bambrick, J., et al. (2024). Accurate structure prediction of biomolecular interactions with alphafold 3. *Nature* 630, 493–500.
43. Jumper, J., Evans, R., Pritzel, A., Green, T., Figurnov, M., Ronneberger, O., Tunyasuvunakool, K., Bates, R., Židek, A., Potapenko, A., et al. (2021). Highly accurate protein structure prediction with alphafold. *Nature* 596, 583–589. <https://doi.org/10.1038/s41586-021-03819-2>.
44. Romera-Paredes, B., Barekatain, M., Novikov, A., Balog, M., Kumar, M. P., Dupont, E., Ruiz, F.J.R., Ellenberg, J.S., Wang, P., Fawzi, O., et al. (2024). Mathematical discoveries from program search with large language models. *Nature* 625, 468–475. <https://doi.org/10.1038/s41586-023-06924-6>.
45. Lam, R., Sanchez-Gonzalez, A., Willson, M., Wirmsberger, P., Fortunato, M., Alet, F., Ravuri, S., Ewalds, T., Eaton-Rosen, Z., Hu, W., et al. (2023). Learning skillful medium-range global weather forecasting. *Science* 382, 1416–1421. <https://doi.org/10.1126/science.adf2336>.
46. Merchant, A., Batzner, S., Schoenholz, S.S., Aykol, M., Cheon, G., and Cubuk, E.D. (2023). Scaling deep learning for materials discovery. *Nature* 624, 80–85. <https://doi.org/10.1038/s41586-023-06735-9>.
47. Bommasani, R., Hudson, D.A., Adeli, E., Altman, R., Arora, S., Arx, S.v., Bernstein, M.S., Bohg, J., Bosselut, A., Brunskill, E., et al. (2021). On the opportunities and risks of foundation models. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2108.07258>.
48. Xu, F., Hao, Q., Zong, Z., Wang, J., Zhang, Y., Wang, J., Lan, X., Gong, J., Ouyang, T., Meng, F., et al. (2025). Towards large reasoning models: A survey of reinforced reasoning with large language models. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2501.09686>.
49. Bahri, Y., Dyer, E., Kaplan, J., Lee, J., and Sharma, U. (2024). Explaining neural scaling laws. e2311878121. *Proc. Natl. Acad. Sci. USA* 121. <https://doi.org/10.1073/pnas.2311878121>.
50. Henighan, T., Kaplan, J., Katz, M., Chen, M., Hesse, C., Jackson, J., Jun, H., Brown, T.B., Dhariwal, P., Gray, S., et al. (2020). Scaling laws for autoregressive generative modeling. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2010.14701>.
51. Hoffmann, J., Borgeaud, S., Mensch, A., Buchatskaya, E., Cai, T., Rutherford, E., de Las Casas, D., Hendricks, L.A., Welbl, J., Clark, A., et al. (2022). Training compute-optimal large language models. In *Proceedings of the 36th International Conference on Neural Information Processing Systems. NIPS '22 (Curran Associates Inc.)*, pp. 30016–30030.
52. Kaplan, J., McCandlish, S., Henighan, T., Brown, T.B., Chess, B., Child, R., Gray, S., Radford, A., Wu, J., and Amodei, D. (2020). Scaling laws for neural language models. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2001.08361>.
53. Ganguli, D., Hernandez, D., Lovitt, L., Askell, A., Bai, Y., Chen, A., Conerly, T., Dassarma, N., Drain, D., Elhage, N., et al. (2022). Predictability and surprise in large generative models. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency. FAccT '22 (Association for Computing Machinery)*, pp. 1747–1764. <https://doi.org/10.1145/3531146.3533229>.
54. Ruan, Y., Maddison, C.J., and Hashimoto, T. (2024). Observational scaling laws and the predictability of language model performance. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*.
55. Schaeffer, R., Schoelkopf, H., Miranda, B., Mukobi, G., Madan, V., Ibrahim, A., Bradley, H., Biderman, S., and Koyejo, S. (2024). Why has predicting downstream capabilities of frontier AI models with scale remained elusive?. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2406.04391>.
56. Brown, T.B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. (2020). Language models are few-shot learners. In *Proceedings of the 34th International Conference on Neural Information Processing Systems. NIPS '20 (Curran Associates Inc.)*, pp. 1877–1901.
57. OpenAI (2024). Learning to reason with LLMs. <https://openai.com/index/learning-to-reason-with-llms/>.
58. Guo, D., Yang, D., Zhang, H., Song, J., Zhang, R., Xu, R., Zhu, Q., Ma, S., Wang, P., Bi, X., et al. (2025). Deepseek-r1: Incentivizing reasoning capability in LLMs via reinforcement learning. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2501.12948>.
59. Brown, B., Juravsky, J., Ehrlich, R., Clark, R., Le, Q.V., Ré, C., and Mirhoseini, A. (2024). Large language monkeys: Scaling inference compute with repeated sampling. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2407.21787>.
60. Snell, C., Lee, J., Xu, K., and Kumar, A. (2024). Scaling LLM test-time compute optimally can be more effective than scaling model parameters. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2408.03314>.
61. Stroebl, B., Kapoor, S., and Narayanan, A. (2024). Inference scaling flaws: The limits of LLM resampling with imperfect verifiers. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2411.17501>.
62. Wu, Y., Sun, Z., Li, S., Welleck, S., and Yang, Y. (2024). Inference scaling laws: An empirical analysis of compute-optimal inference for problem-solving with language models. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2408.00724>.
63. Li, M., Kudugunta, S., and Zettlemoyer, L. (2025). (mis)fitting scaling laws: A survey of scaling law fitting techniques in deep learning. In *The Thirteenth International Conference on Learning Representations*.
64. Anderson, P.W. (1972). More is different. *Science* 177, 393–396. <https://doi.org/10.1126/science.177.4047.393>.
65. Stanley, H.E. (1987). Introduction to phase transitions and critical phenomena. In *International Series of Monographs on Physics (Oxford university press)*.

66. Lubana, E.S., Kawaguchi, K., Dick, R.P., and Tanaka, H. (2024). A percolation model of emergence: Analyzing transformers trained on a formal language. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2408.12578>.
67. Pan, A., Bhatia, K., and Steinhart, J. (2021). The effects of reward mis-specification: Mapping and mitigating misaligned models. In *International Conference on Learning Representations*.
68. Epoch AI (2024). Frontiermath. <https://epoch.ai/frontiermath>.
69. Fu, R., Jin, G.Z., and Liu, M. (2022). Does human-algorithm feedback loop lead to error propagation? Evidence from zillow's zestimate. In *National Bureau of Economic Research (National Bureau of Economic Research)*. <https://doi.org/10.3386/w29880>.
70. Hardt, M., and Mendler-Düner, C. (2023). Performative prediction: Past and future. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2310.16608>.
71. Healy, K. (2015). The performativity of networks. *Arch. Eur. Sociol.* 56, 175–205. <https://doi.org/10.1017/S0003975615000107>.
72. Perdomo, J., Zmc, T., Mendler-Düner, C., and Hardt, M. (2020). Performative prediction. In *Proceedings of the 37th International Conference on Machine Learning (PMLR)*, pp. 7599–7609. <https://proceedings.mlr.press/v119/perdomo20a.html>.
73. Cen, S.H., Ilyas, A., Allen, J., Li, H., and Madry, A. (2024). Measuring stratification in recommendation: Users adapt their behavior to shape future content. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2405.05596>.
74. Pedreschi, D., Pappalardo, L., Ferragina, E., Baeza-Yates, R., Barabási, A.L., Dignum, F., Dignum, V., Eliassi-Rad, T., Giannotti, F., Kertész, J., et al. (2025). Human-AI coevolution. *Artif. Intell.* 339, 104244. <https://doi.org/10.1016/j.artint.2024.104244>.
75. Stray, J., Halevy, A., Assar, P., Hadfield-Menell, D., Boutilier, C., Ashar, A., Bakalar, C., Beattie, L., Ekstrand, M., Leibowicz, C., et al. (2024). Building human values into recommender systems: An interdisciplinary synthesis. *ACM Trans. Recomm. Syst.* 2, 1–20. <https://doi.org/10.1145/3632297>.
76. Williams, M., Carroll, M., Narang, A., Weisser, C., Murphy, B., and Dragan, A. (2024). On targeted manipulation and deception when optimizing LLMs for user feedback. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2411.02306>.
77. Bender, E.M., Gebru, T., McMillan-Major, A., and Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency. FAccT '21 (Association for Computing Machinery)*, pp. 610–623. <https://doi.org/10.1145/3442188.3445922>.
78. Shur-Ofry, M. (2025). Multiplicity as an AI governance principle. *Indiana Law Journal* 100, 6. <https://www.repository.law.indiana.edu/ilj/vol100/iss4/6/>.
79. Taori, R., and Hashimoto, T.B. (2023). Data feedback loops: model-driven amplification of dataset biases. In *Proceedings of the 40th International Conference on Machine Learning vol. 202 of ICML'23 (JMLR)*, pp. 33883–33920.
80. Kolt, N. (2022). Predicting consumer contracts. *Berk. Technol. Law J.* 37, 71–138.
81. Veselovsky, V., Ribeiro, M.H., Cozzolino, P., Gordon, A., Rothschild, D., and West, R. (2023). Prevalence and prevention of large language model use in crowd work. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2310.15683>.
82. Veselovsky, V., Ribeiro, M.H., and West, R. (2023). Artificial artificial intelligence: Crowd workers widely use large language models for text production tasks. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2306.07899>.
83. Wu, T., Zhu, H., Albayrak, M., Axon, A., Bertsch, A., Deng, W., Ding, Z., Guo, B., Gururaja, S., Kuo, T.S., et al. (2023). LLMs as workers in human-computational algorithms? replicating crowdsourcing pipelines with llms. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2307.10168>.
84. Jakesch, M., Bhat, A., Buschek, D., Zalmanson, L., and Naaman, M. (2023). Co-writing with opinionated language models affects users' views. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems. CHI '23 (Association for Computing Machinery)*, pp. 1–15. <https://doi.org/10.1145/3544548.3581196>.
85. Fan, L., Chen, K., Krishnan, D., Katabi, D., Isola, P., and Tian, Y. (2024). Scaling laws of synthetic images for model training... for now. In *2024 IEEE-CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 7382–7392. <https://doi.org/10.1109/CVPR52733.2024.00705>.
86. Lee, H., Phatale, S., Mansoor, H., Mesnard, T., Ferret, J., Lu, K., Bishop, C., Hall, E., Carbune, V., Rastogi, A., and Prakash, S. (2024). RLHF vs. RLHF: Scaling reinforcement learning from human feedback with AI feedback. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2309.00267>.
87. Liu, R., Wei, J., Liu, F., Si, C., Zhang, Y., Rao, J., Zheng, S., Peng, D., Yang, D., Zhou, D., and Dai, A.M. (2024). Best practices and lessons learned on synthetic data. In *First Conference on Language Modeling*. <https://openreview.net/forum?id=QJaWBh61C>.
88. Singh, A., Co-Reyes, J.D., Agarwal, R., Anand, A., Patil, P., Garcia, X., Liu, P.J., Harrison, J., Lee, J., Xu, K., et al. (2024). Beyond human data: Scaling self-training for problem-solving with language models. In *Transactions on Machine Learning Research*.
89. Alemohammad, S., Casco-Rodríguez, J., Luzi, L., Humayun, A.I., Babaei, H., LeJeune, D., Siahkoobi, A., and Baraniuk, R. (2023). Self-consuming generative models go mad. In *The Twelfth International Conference on Learning Representations*.
90. Bohacek, M., and Farid, H. (2023). Nepotistically trained generative-AI models collapse. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2311.12202>.
91. Gerstgrasser, M., Schaeffer, R., Dey, A., Rafailov, R., Korbak, T., Sleight, H., Agrawal, R., Hughes, J., Pai, D.B., Gromov, A., et al. (2024). Is model collapse inevitable? breaking the curse of recursion by accumulating real and synthetic data. In *First Conference on Language Modeling*. <https://openreview.net/forum?id=5B2K4LRgmz>.
92. Kazdan, J., Schaeffer, R., Dey, A., Gerstgrasser, M., Rafailov, R., Donoho, D.L., and Koyejo, S. (2024). Collapse or thrive? perils and promises of synthetic data in a self-generating world. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2410.16713>.
93. Shumailov, I., Shumaylov, Z., Zhao, Y., Gal, Y., Papernot, N., and Anderson, R. (2024). The curse of recursion: Training on generated data makes models forget. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2305.17493>.
94. Shumailov, I., Shumaylov, Z., Zhao, Y., Papernot, N., Anderson, R., and Gal, Y. (2024). AI models collapse when trained on recursively generated data. *Nature* 631, 755–759. <https://doi.org/10.1038/s41586-024-07566-y>.
95. Kirchenbauer, J., Geiping, J., Wen, Y., Katz, J., Miers, I., and Goldstein, T. (2023). A watermark for large language models. In *Proceedings of the 40th International Conference on Machine Learning (PMLR)*, pp. 17061–17084. <https://proceedings.mlr.press/v202/kirchenbauer23a.html>.
96. Kirchenbauer, J., Geiping, J., Wen, Y., Shu, M., Saifullah, K., Kong, K., Fernando, K., Saha, A., Goldblum, M., and Goldstein, T. (2023). On the reliability of watermarks for large language models. In *The Twelfth International Conference on Learning Representations*.
97. Sadasivan, V.S., Kumar, A., Balasubramanian, S., Wang, W., and Feizi, S. (2024). Can ai-generated text be reliably detected?. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2303.11156>.
98. Huang, S., and Siddarth, D. (2023). Generative AI and the digital commons. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2303.11074>.
99. Liang, W., Izzo, Z., Zhang, Y., Lepp, H., Cao, H., Zhao, X., Chen, L., Ye, H., Liu, S., Huang, Z., et al. (2024). Monitoring ai-modified content at scale: A case study on the impact of chatgpt on AI conference peer reviews. In *Forty-first International Conference on Machine Learning*.
100. Thompson, B., Dhaliwal, M., Frisch, P., Domhan, T., and Federico, M. (2024). A shocking amount of the web is machine translated: Insights from multi-way parallelism. In *Findings of the Association for Computational Linguistics: ACL 2024*, L.W. Ku, A. Martins, and V. Srikumar, eds. (Association for Computational Linguistics), pp. 1763–1775. <https://doi.org/10.18653/v1/2024.findings-acl.103>.
101. Balloccu, S., Schmidová, P., Lango, M., and Dusek, O. (2024). Leak, cheat, repeat: Data contamination and evaluation malpractices in

- closed-source LLMs. In Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers), Y. Graham and M. Purver, eds. (Association for Computational Linguistics), pp. 67–93. <https://aclanthology.org/2024.eacl-long.5>.
102. Panickssery, A., Bowman, S.R., and Feng, S. (2024). LLM evaluators recognize and favor their own generations. In The Thirty-eighth Annual Conference on Neural Information Processing Systems. <https://openreview.net/forum?id=4NJBV6Wp0h>.
103. Perez, E., Huang, S., Song, F., Cai, T., Ring, R., Aslanides, J., Glaese, A., McAleese, N., and Irving, G. (2022). Red Teaming Language Models with Language Models. In Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, Y. Goldberg, Z. Kozareva, and Y. Zhang, eds. (Association for Computational Linguistics), pp. 3419–3448. <https://doi.org/10.18653/v1/2022.emnlp-main.225>.
104. Perez, E., Ringer, S., Lukosiute, K., Nguyen, K., Chen, E., Heiner, S., Pettit, C., Olsson, C., Kundu, S., Kadavath, S., et al. (2023). Discovering language model behaviors with model-written evaluations. In Findings of the Association for Computational Linguistics: ACL 2023, A. Rogers, J. Boyd-Graber, and N. Okazaki, eds. (Association for Computational Linguistics), pp. 13387–13434. <https://doi.org/10.18653/v1/2023.findings-acl.847>.
105. Bowman, S.R., Hyun, J., Perez, E., Chen, E., Pettit, C., Heiner, S., Lukosiute, K., Askell, A., Jones, A., Chen, A., et al. (2022). Measuring Progress on Scalable Oversight for Large Language Models. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2211.03540>.
106. Burns, C., Izmailov, P., Kirchner, J.H., Baker, B., Gao, L., Aschenbrenner, L., Chen, Y., Ecoffet, A., Joglekar, M., Leike, J., et al. (2024). Weak-to-strong generalization: Eliciting strong capabilities with weak supervision. In Forty-first International Conference on Machine Learning.
107. Pan, A., Jones, E., Jagadeesan, M., and Steinhardt, J. (2024). Feedback loops with language models drive in-context reward hacking. In Forty-first International Conference on Machine Learning.
108. Wyllie, S., Shumailov, I., and Papernot, N. (2024). Training on synthetic data amplifies bias. In Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency. FAccT '24 (Association for Computing Machinery), pp. 2113–2147. Fairness feedback loops. <https://doi.org/10.1145/3630106.3659029>.
109. Feng, S., Park, C.Y., Liu, Y., and Tsvetkov, Y. (2023). From Pretraining Data to Language Models to Downstream Tasks: Tracking the Trails of Political Biases Leading to Unfair Nlp Models. In Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), A. Rogers, J. Boyd-Graber, and N. Okazaki, eds. (Association for Computational Linguistics), pp. 11737–11762. <https://doi.org/10.18653/v1/2023.acl-long.656>.
110. Toups, C., Bommasani, R., Creel, K.A., Bana, S.H., Jurafsky, D., and Liang, P. (2024). Ecosystem-level analysis of deployed machine learning reveals homogeneous outcomes. In Proceedings of the 37th International Conference on Neural Information Processing Systems. NIPS '23 (Curran Associates Inc.), pp. 51178–51201.
111. Betley, J., Tan, D., Warncke, N., Szyber-Betley, A., Bao, X., Soto, M., Labenz, N., and Evans, O. (2025). Emergent misalignment: Narrow finetuning can produce broadly misaligned LLMs. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2502.17424>.
112. Zou, A., Wang, Z., Carlini, N., Nasr, M., Kolter, J.Z., and Fredrikson, M. (2023). Universal and transferable adversarial attacks on aligned language models. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2307.15043>.
113. Li, A., Zhou, Y., Raghuram, V.C., Goldstein, T., and Goldblum, M. (2025). Commercial LLM agents are already vulnerable to simple yet dangerous attacks. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2502.08586>.
114. Chan, A., Salganik, R., Markelius, A., Pang, C., Rajkumar, N., Krasheninnikov, D., Langosco, L., He, Z., Duan, Y., Carroll, M., et al. (2023). Harms from increasingly agentic algorithmic systems. In Proceedings of the 2023 ACM Conference on Fairness, Accountability, and Transparency. FAccT '23 (Association for Computing Machinery), pp. 651–666. <https://doi.org/10.1145/3593013.3594033>.
115. Cohen, M.K., Kolt, N., Bengio, Y., Hadfield, G.K., and Russell, S. (2024). Regulating advanced artificial agents. *Science* 384, 36–38. <https://doi.org/10.1126/science.adl0625>.
116. Kolt, N. (2025). Governing AI agents. *Notre Dame Law Review* (forthcoming) 101. <https://papers.ssrn.com/abstract%3D4772956>
117. Ruan, Y., Dong, H., Wang, A., Pitis, S., Zhou, Y., Ba, J., Dubois, Y., Madison, C.J., and Hashimoto, T. (2024). Identifying the risks of lm agents with an lm-emulated sandbox. In The Twelfth International Conference on Learning Representations.
118. Shur-Ofry, M. (2024). A networks-of-networks perspective on AI policy. In *Network Law Review*. <https://www.networklawreview.org/shur-ofry-generative-ai/>.
119. Hammond, L., Chan, A., Clifton, J., Hoelscher-Obermaier, J., Khan, A., McLean, E., Smith, C., Barfuss, W., Foerster, J., Gavenciak, T., et al. (2025). Multi-agent risks from advanced ai. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2502.14143>.
120. Taylor, J.B., and Williams, J.C. (2009). A black swan in the money market. *American Economic Journal: Macroeconomics* 1, 58–83. <https://doi.org/10.1257/mac.1.1.58>.
121. Richards, C.E., Tzachor, A., Avin, S., and Fenner, R. (2023). Rewards, risks and responsible deployment of artificial intelligence in water systems. *Nat. Water* 1, 422–432. <https://doi.org/10.1038/s44221-023-00069-6>.
122. Galaz, V., Centeno, M.A., Callahan, P.W., Causevic, A., Patterson, T., Brass, I., Baum, S., Farber, D., Fischer, J., Garcia, D., et al. (2021). Artificial intelligence, systemic risks, and sustainability. *Technol. Soc.* 67, 101741. <https://doi.org/10.1016/j.techsoc.2021.101741>.
123. Tzachor, A., Devare, M., King, B., Avin, S., and Ó hÉigeartaigh, S. (2022). Responsible artificial intelligence in agriculture requires systemic understanding of risks and externalities. *Ó hÉigeartaigh, S. Nat. Mach. Intell.* 4, 104–109.
124. Askell, A., Brundage, M., and Hadfield, G. (2019). The role of cooperation in responsible AI development. Preprint at arXiv. <https://doi.org/10.48550/arXiv.1907.04534>.
125. Kolt, N. (2024). Algorithmic Black Swans. *Washington University Law Review* 101, 1177–1240.
126. Dafoe, A. (2023). Overview and theoretical lenses. In *The Oxford Handbook of AI Governance*, J.B. Bullock, Y.C. Chen, J. Himmelreich, V.M. Hudson, A. Korinek, M.M. Young, and B. Zhang, eds. (Oxford University Press), pp. 21–44. AI governance. <https://doi.org/10.1093/oxfordhb/9780197579329.013.2>.
127. Arbel, Y., Tokson, M., and Lin, A. (2024). Systemic regulation of artificial intelligence. *Ariz. State Law J.* 56, 545.
128. Centeno, M.A., Nag, M., Patterson, T.S., Shaver, A., and Windawi, A.J. (2015). The emergence of global systemic risk. *Annu. Rev. Sociol.* 41, 65–85.
129. Renn, O., Lucas, K., Haas, A., and Jaeger, C. (2019). Things are different today: the challenge of global systemic risks. *J. Risk Res.* 22, 401–415.
130. Schweizer, P.J., and Juhola, S. (2024). Navigating systemic risks: governance of and for systemic risks. *Glob. Sustain.* 7, e38.
131. Arnscheidt, C.W., Beard, S., Hobson, T., Ingram, P., Kemp, L., Mani, L., Marcoci, A., Mbeva, K., hÉigeartaigh, S., Sandberg, A., et al. (2025). Systemic contributions to global catastrophic risk. In *Global Sustainability*. forthcoming. <https://doi.org/10.17863/CAM.117342>.
132. Cosens, B., Ruhl, J.B., Soininen, N., Gunderson, L., Belinskij, A., Blenckner, T., Camacho, A.E., Chaffin, B.C., Craig, R.K., Doremus, H., et al. (2021). Governing complexity: Integrating science, governance, and law to manage accelerating change in the globalized commons. *e2102798118. Proc. Natl. Acad. Sci. USA* 118. <https://doi.org/10.1073/pnas.2102798118>.
133. Craig, R.K. (2010). Stationarity is dead - long live transformation: Five principles for climate change adaptation law. *Harv. Environ. Law Rev.* 34, 9–74.

134. Ruhl, J.B. (2010). General design principles for resilience and adaptive capacity in legal systems - with applications to climate change adaptation and resiliency in legal systems. *N. C. Law Rev.* 89, 1373–1404.
135. Arner, D.W. (2010). Adaptation and resilience in global financial regulation adaptation and resiliency in legal systems. *N. C. Law Rev.* 89, 1579–1628.
136. Schwarcz, S.L. (2008). Systemic risk. *Georgetown Law J.* 97, 193–250.
137. Schwarcz, S.L. (2009). Regulating Complexity in Financial Markets 87 (*Washington University Law Review*), pp. 211–268.
138. Malcai, O., and Shur-Ofry, M. (2021). Using complexity to calibrate legal response to covid-19. *Front. Phys.* 9, 650943. <https://doi.org/10.3389/fphy.2021.650943>.
139. Collingridge, D. (1980). *The social control of technology* (St. Martin's Press).
140. Moses, L.B. (2007). Recurring dilemmas: The law's race to keep up with technological change. *University of Illinois Journal of Law, Technology & Policy* 2007, 239–286.
141. G.E. Marchant, B.R. Allenby, and J.R. Herkert, eds. (2011). *The Growing Gap Between Emerging Technologies and Legal-Ethical Oversight: The Pacing Problem*. The International Library of Ethics, Law and Technology, 7 (Springer Netherlands). <https://doi.org/10.1007/978-94-007-1356-7>.
142. Crotoof, R., and Ard, B.J. (2020). Structuring techlaw. *Harv. J. Law Technol.* 34, 347–418.
143. Kaminski, M.E. (2023). Regulating the Risks of Ai103 (*Boston University Law Review*), pp. 1347–1411.
144. Doremus, H. (2010). Adaptive management as an information problem adaptation and resiliency in legal systems. *N. C. Law Rev.* 89, 1455–1498.
145. Karkkainen, B.C. (2000). Information as environmental regulation: Tri and performance benchmarking, precursor to a new paradigm. *Georgetown Law J.* 89, 257–370.
146. Karkkainen, B.C. (2007). Bottlenecks and baselines: Tackling information deficits in environmental regulation. *Tex. Law Rev.* 86, 1409–1444.
147. Kay, J.A., and King, M.A. (2020). *Radical uncertainty* (The Bridge Street Press).
148. V.A.W.J. Marchau, W.E. Walker, P.J.T.M. Bloemen, and S.W. Popper, eds. (2019). *Decision Making under Deep Uncertainty: From Theory to Practice* (Springer International Publishing). <https://doi.org/10.1007/978-3-030-05252-2>.
149. Ford, C. (2013). Prospects for scalability: Relationships and uncertainty in responsive regulation. *Regulation & Governance* 7, 14–29. <https://doi.org/10.1111/j.1748-5991.2012.01166.x>.
150. Crotoof, R., Kaminski, M.E., and Price, W.N.I. (2023). Humans in the loop. *Vanderbilt Law Rev.* 76, 429–510.
151. Benthall, S., and Sivan-Sevilla, I. (2024). Regulatory ci: Adaptively regulating privacy as contextual integrity. In *ACM Symposium on Computer Science & Law*.
152. Reuel, A., and Undheim, T.A. (2024). Generative AI needs adaptive governance. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2406.04554>.
153. Folke, C., Hahn, T., Olsson, P., and Norberg, J. (2005). Adaptive governance of social-ecological systems. *Annu. Rev. Environ. Resour.* 30, 441–473. <https://doi.org/10.1146/annurev.energy.30.050504.144511>.
154. Holling, C.S. (1973). Resilience and stability of ecological systems. *Annu. Rev. Ecol. Syst.* 4, 1–23.
155. Holling, C.S. (1978). Adaptive environmental assessment and management. In *International Series on Applied Systems Analysis* (Wiley).
156. Ruhl, J.B. (2005). Regulation by adaptive management—is it possible? *Minnesota Journal of Law. Sci. Technol.* 7, 21–58.
157. Duit, A., and Galaz, V. (2008). Governance and complexity—emerging issues for governance theory. *Governance* 21, 311–335.
158. D'Souza, R.M. (2017). Curtailing cascading failures. *Science* 358, 860–861. <https://doi.org/10.1126/science.aag0474>.
159. Ruhl, J.B. (2019). Governing cascade failures in complex social-ecological-technological systems: Framing context, strategies, and challenges. *Vanderbilt Journal of Entertainment & Technology Law* 22, 407–440.
160. Fisher, L., and Sandberg, A. (2022). A safe governance space for humanity: necessary conditions for the governance of global catastrophic risks. *Glob. Policy* 13, 792–807.
161. Vivo, P., Katz, D.M., and Ruhl, J.B. (2024). A complexity science approach to law and governance. *Philosophical Transactions of the Royal Society A: Mathematical, Phys. Eng. Sci.* 382, 20230166. <https://doi.org/10.1098/rsta.2023.0166>.
162. National Institute of Standards and Technology (2023). AI risk management framework. NIST. <https://www.nist.gov/itl/ai-risk-management-framework>.
163. People's Republic of China (2023). Interim measures for the management of generative artificial intelligence services. https://www.cac.gov.cn/2023-07/13/c_1690898327029107.htm.
164. People's Republic of China (2023). Artificial intelligence law of the people's republic of china. <https://perma.cc/L9E4-5K3V>.
165. White House (2023). Executive order on the safe, secure, and trustworthy development and use of artificial intelligence.
166. White House (2025). Removing barriers to american leadership in artificial intelligence. <https://www.whitehouse.gov/presidential-actions/2025/01/removing-barriers-to-american-leadership-in-artificial-intelligence/>.
167. European Union (2025). General-purpose AI code of practice. <https://digital-strategy.ec.europa.eu/en/policies/ai-code-practice>.
168. Coglianese, C., Zeckhauser, R., and Parson, E. (2004). Seeking truth for power: Informational strategy and regulatory policymaking. *Minn. Law Rev.* 89, 277–341.
169. Stephenson, M.C. (2010). Information acquisition and institutional design. *Harv. Law Rev.* 124, 1422–1484.
170. Van Loo, R. (2019). Regulatory monitors: Policing firms in the compliance era. *Columbia Law Rev.* 119, 369–444.
171. Whittlestone, J., and Clark, J. (2021). Why and how governments should monitor AI development. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2108.12427>.
172. Clark, J. (2023). Information markets and AI development. In *The Oxford Handbook of AI Governance*, J.B. Bullock, Y.C. Chen, J. Himmelreich, V.M. Hudson, A. Korinek, M.M. Young, and B. Zhang, eds. (Oxford University Press), pp. 345–357. <https://doi.org/10.1093/oxfordhb/9780197579329.013.21>.
173. Kolt, N., Anderljung, M., Barnhart, J., Brass, A., Esvelt, K., Hadfield, G.K., Heim, L., Rodriguez, M., Sandbrink, J.B., Woodside, T., et al. (2024). Responsible reporting for frontier AI development. Proceedings of the AAAI-ACM Conference on AI, Ethics, and Society 7, 768–783. <https://doi.org/10.1609/aies.v7i1.31678>.
174. Reuel, A., Bucknall, B., Casper, S., Fist, T., Soder, L., Aarne, O., Hammond, L., Ibrahim, L., Chan, A., Wills, P., et al. (2024). Open problems in technical AI governance. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2407.14981>.
175. Casper, S., Krueger, D., and Hadfield-Menell, D. (2025). Pitfalls of evidence-based AI policy. In *International Conference on Learning Representations*.
176. Oreskes, N., and Conway, E.M. (2010). *Merchants of doubt: How a handful of scientists obscured the truth on issues from tobacco smoke to global warming* (Bloomsbury Publishing).
177. Koessler, L., Schuett, J., and Anderljung, M. (2024). Risk thresholds for frontier ai. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2406.14713>.
178. Galle, B. (2015). In praise of ex ante regulation. *Vanderbilt Law Rev.* 68, 1715–1760.

179. Posner, R.A. (2004). *Catastrophe: Risk and Response* (Oxford University Press). <https://academic.oup.com/book/40836>.
180. Sunstein, C.R. (2005). Irreversible and catastrophic. *Cornell Law Rev.* 91, 841–898.
181. Simon, H.A. (1955). A behavioral model of rational choice. *Q. J. Econ.* 69, 99–118. <https://doi.org/10.2307/1884852>.
182. Harremoës, P., Gee, D., MacGarvin, M., Stirling, A., Keys, J., Wynne, B., and Vaz, S.G. (2013). *The Precautionary Principle in the 20th Century: Late Lessons from Early Warnings* (Taylor and Francis).
183. Nash, J.R. (2008). Standing and the precautionary principle. *Columbia Law Rev.* 108, 494–528.
184. Tegmark, M., and Omohundro, S. (2023). Provably safe systems: the only path to controllable agi. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2309.01933>.
185. Dalrymple, D., Skalse, J., Bengio, Y., Russell, S., Tegmark, M., Seshia, S., Omohundro, S., Szegedy, C., Goldhaber, B., Ammann, N., et al. (2024). Towards guaranteed safe AI: A framework for ensuring robust and reliable AI systems. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2405.06624>.
186. Buhl, M.D., Sett, G., Koessler, L., Schuett, J., and Anderljung, M. (2024). Safety cases for frontier AI. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2410.21572>.
187. Clymer, J., Gabrieli, N., Krueger, D., and Larsen, T. (2024). Safety cases: How to justify the safety of advanced AI systems. Preprint at arXiv. <https://doi.org/10.48550/arXiv.2403.10462>.
188. Farber, D.A. (2010). Uncertainty. *Georgetown Law J.* 99, 901–960.
189. Sunstein, C.R. (2003). Beyond the precautionary principle. *Univ. PA. Law Rev.* 151, 1003–1058.
190. Simon, H.A. (1962). The architecture of complexity. *Proc. Am. Phil. Soc.* 106, 467–482.

About the authors

Noam Kolt is an assistant professor at the Hebrew University Faculty of Law and School of Computer Science and Engineering. He leads the Governance of AI Lab (GOAL), a cross-disciplinary research group developing technical and institutional infrastructure to support safe and ethical AI. During his doctorate at the University of Toronto, Noam served as a research advisor to Google DeepMind and was a member of OpenAI's GPT-4 red team. He has published in the *Washington University Law Review*, *Berkeley Technology Law Journal*, *Yale Law & Policy Review*, and peer-reviewed venues, including *NeurIPS*, *ACM FAccT*, and *Science*.

Michal Shur-Ofry is an associate professor at the Hebrew University Law Faculty and an affiliate visiting faculty at NYU Information Law Institute. A graduate of Hebrew University (LLB, PhD) and University College London (LLM), she held visiting and teaching positions at Georgetown University, the University of British Columbia, and NYU. Her research bridges law and complex systems science, using insights and formal tools from complexity to inform legal rules and policies across diverse legal domains. She also examines regulatory responses to systemic implications of AI. Her research has been published in leading journals in law, public policy, physics, and metascience.

Reuven Cohen is a professor of mathematics at Bar-Ilan University. He is a graduate of Bar-Ilan University (BSc in Physics and Computer Science and PhD in Physics) and has held postdoctoral research positions at the Weizmann Institute of Science, Boston University, and MIT. His research focuses on the theory of complex networks, geometric and combinatorial optimization algorithms, and applied mathematics in general. He has published research in leading journals in physics, computer science, engineering, and mathematics, as well as in interdisciplinary research in law and policy.